

# Test #1

截止日期 5月5日 23:59

分數 100

問題 40

可用 5月1日 0:01 - 5月5日 23:59

時間限制 60 分鐘

## 說明

**Canvas calls this a "Quiz", but it is really Test #1.**

**It consists of 40 multiple choice questions to be done in 60 minutes. It is Open Notes.**

**Once you start, you must finish. Canvas will not let you pause and come back.**

## 嘗試記錄

	嘗試	時間	分數
最新的	<a href="#">嘗試 1</a>	43 分鐘	得分：100；總分：100

❗ 正確答案將於 5月6日 0:01 可用。

此測驗的分數： 得分：100；總分：100

已提交5月3日 12:15

此嘗試持續 43 分鐘。



問題 1

2.5 / 2.5 分數

**Our class's "Inverse Amdahl's Law" that you used in Projects #0 and #1 computes:**

- ☒  $F_p$ , given  $S_n$  and  $n$
- ☐  $S_n$ , given  $F_p$  and  $n$
- ☐ Thread Efficiency, given  $S_n$  and  $n$
- ☐  $n$ , given  $S_n$  and  $F_p$



## 問題 2

2.5 / 2.5 分數

**In terms of 8-byte double-precision numbers, the size of a cache line is:**

- ☒ 8 double-precision numbers
- ☐ 4 double-precision numbers
- ☐ 16 double-precision numbers
- ☐ 32 double-precision numbers



## 問題 3

2.5 / 2.5 分數

**The difference between static and dynamic scheduling of an OpenMP for-loop is:**

- ☐ Dynamic scheduling divides all the for-loop passes among the threads at first
- ☒ Dynamic scheduling divides only some of the for-loop passes among the threads at first
- ☐ Dynamic scheduling allows you to change how the for-loop passes are divided up while they are running
- ☐ Dynamic scheduling changes the chunksize while the for-loop is running



## 問題 4

2.5 / 2.5 分數

**The two types of coherence that caches want to see in order to deliver maximum performance are:**

- ☒ Spatial and Temporal
- ☐ Spatial and Thermal
- ☐ Systemic and Temporal
- ☐ Systemic and Thermal



## 問題 5

2.5 / 2.5 分數

**A Deadlock condition is when:**

- ☒ Two threads are each waiting for the other one to do something
- ☐ The CPU chip cannot find any more instructions to execute while waiting for a memory fetch
- ☐ When you keep internal state
- ☐ When it is a race to see which of two threads get to a piece of code first



## 問題 6

2.5 / 2.5 分數

**Hyperthreading is:**

- ☐ Adding extra cache space
- ☒ Keeping one or more extra thread states within a core
- ☐ Adding more memory bandwidth
- ☐ Adding one or more cores



## 問題 7

2.5 / 2.5 分數

**A good way to make a piece of code *not* Thread Safe is to:**

- ☐ Use a mutual exclusion lock
- ☐ Use a private variable
- ☒ Keep internal state
- ☐ Use a chunksize of 1



## 問題 8

2.5 / 2.5 分數

**A Private variable differs from a Shared variable in that:**

- ☐ When each thread writes to it, the value goes to the same memory address
- ☒ Each thread has its own copy of it
- ☐ Writing to it automatically triggers a cache line reload
- ☐ Writing to it automatically triggers a power-of-two reduction operation



## 問題 9

2.5 / 2.5 分數

**OpenMP Reductions are faster than Atomic or Critical because:**

- ☐ They sum into an array whose elements are a Fibonacci series in size
- ☐ They momentarily disable interrupts to keep the summing equation from being corrupted
- ☐ They sum into a user-supplied array and then let the programmer decide how to best sum them
- ☒ They sum into a separate variable per thread and then perform power-of-two addition



## 問題 10

2.5 / 2.5 分數

**To get an A in CS 475/575 requires:**

- ☐ A weighted average of 93%
- ☒ 1060 total points
- ☐ A weighted average of 96%



## 問題 11

2.5 / 2.5 分數

**The word “deterministic” means:**

- ☒ The same inputs will always produce the same outputs
- ☐ The program outputs change whenever you change the number of threads
- ☐ It describes a quantity that you are attempting to determine
- ☐ The program outputs change every time you run the program



## 問題 12

2.5 / 2.5 分數

**Why is there a photo of a carton of eggs in the Cache notes?**

- ☐ It explains Temporary Coherence
- ☒ Bringing home a dozen eggs when you only need 2 today is like the way cache works
- ☐ Because the size of a cache line is a dozen floats
- ☐ It explains Stationary Coherence



## 問題 13

2.5 / 2.5 分數

**When multiplying two arrays together, you've decided you want to do it **4-pairs-of-numbers** at a time. The type of SIMD you should use is:**

- ☒ SSE
- ☐ AVX
- ☐ AVX-512
- ☐ MMX



## 問題 14

2.5 / 2.5 分數

**A “race condition” is one where:**

- ☒ You get a different result depending on which thread gets to a piece of code first
- ☐ You get the same result regardless of which thread gets to a piece of code first
- ☐ It matters which stack holds a particular variable
- ☐ It matters which thread gets to a barrier first



## 問題 15

2.5 / 2.5 分數

**In an n-core multicore program, what do you need to do to compute the  $F_{\text{parallel}}$ ?**

- ☐ Go find out the size of the cache and use the inverse Amdahl's Law
- ☐ Figure out how many CPU sockets are in use and use the inverse Amdahl's Law
- ☐ Measure just the 20-core performance and use the inverse Amdahl's Law
- ☒ Measure the Speedup and use our inverse Amdahl's Law



## 問題 16

2.5 / 2.5 分數

**Which of these is an example of a forbidden *inter-loop dependency*?**

- ☐ `a[i] = (float)( i );`
- ☒ `a[ i ] = a[ i-1 ] + 1.;`
- ☐ `a[ i ] = b[ i ] + 1.;`
- ☐ `a[ i ] = 2.*a[ i ];`



## 問題 17

2.5 / 2.5 分數

**The reason that our OpenMP programs have a NUMTRIES for-loop is to:**

- ☐ Determine the median performance
- ☒ Determine the peak performance
- ☐ Determine the standard deviation of performance
- ☐ Determine the range of performance numbers



## 問題 18

2.5 / 2.5 分數

**The line “#pragma omp single” is used to:**

- ☐ Force this block of code to undergo a single reduction
- ☐ Force this block of code to be divided up into individual OpenMP sections
- ☒ Force this block of code to be executed by one thread only
- ☐ Force this block of code to be executed in single-file order by each thread



## 問題 19

2.5 / 2.5 分數

**The cache that is smallest and fastest is named:**

- ☐ L2
- ☒ L1
- ☐ L3
- ☐ L0



## 問題 20

2.5 / 2.5 分數

**The cache that is closest to the Arithmetic Logic Unit (ALU) is named:**

- ☐ L0
- ☐ L3
- ☐ L2
- ☒ L1



## 問題 21

2.5 / 2.5 分數

**In CS 475/575, the maximum number of Bonus Days that you can use on any one projects is:**

- ☒ 2
- ☐ 5
- ☐ 4
- ☐ 3



## 問題 22

2.5 / 2.5 分數

**One way to prevent harm from race conditions is:**

- ☐ Dynamic scheduling
- ☐ Private variables
- ☐ Shared variables
- ☒ Mutual Exclusion Locks



## 問題 23

2.5 / 2.5 分數

**Coarse-grained parallelism is:**

- ☐ Dividing the problem into a large number of small pieces
- ☒ Dividing the problem into a small number of large pieces
- ☐ Dividing the problem into equal-size pieces
- ☐ Dividing the problem into pieces, of all which have to be a different size



## 問題 24

2.5 / 2.5 分數

**In multithreading, the threads all share:**

- ☐ Heap, Global variables, and the same Stack
- ☐ Execution instructions, Global variables, and the same Stack
- ☐ Heap, Execution instructions, and the same Stack
- ☒ Heap, Execution instructions, and Global variables



## 問題 25

2.5 / 2.5 分數

**A "Mutex" is:**

- ☐ A "multiple texture" for graphics processing
- ☐ A "mutual text" message
- ☐ A sound you make when you sneeze
- ☒ Another term for a "mutual exclusion lock"



## 問題 26

2.5 / 2.5 分數

**A Barrier is:**

- ☐ A place in the code where the first thread to get there issues an interrupt
- ☒ A place in the code that all threads must reach before any of them are allowed to continue
- ☐ A place in the code that threads are not allowed to pass ever
- ☐ A place in the code where threads can spawn other threads



## 問題 27

2.5 / 2.5 分數

**False Sharing happens because:**

- ☒ One core is writing to a cache line at the same time another core is reading or writing the same cache line
- ☐ Two cores are not sharing the same cache line, but should be
- ☐ One core is writing to a cache line at the same time another core is reading or writing a different cache line
- ☐ Two cores are reading from the same cache line



## 問題 28

2.5 / 2.5 分數

When adding up the elements of a 2D array in C or C++, it is faster to add the elements:

- ☐ Vertically (i.e., down the columns) first
- ☒ Horizontally (i.e., across the rows) first
- ☐ It makes no speed difference either way



## 問題 29

2.5 / 2.5 分數

**SPMD stands for:**

- ☒ Single Program, Multiple Data
- ☐ Significant Parallelism, Multiple Data
- ☐ Single Program, Much Data
- ☐ Significant Parallelism, Much Data



## 問題 30

2.5 / 2.5 分數

**How many total Bonus Days are you allowed in CS 475/575?**



- ☐ 3
- ☐ 2
- ☒ 5
- ☐ 4
- ☐ 6



問題 31

2.5 / 2.5 分數

**The theoretical maximum speedup that you can ever achieve, no matter how many cores you add, is:**

- ☐  $1/(F_p + F_s)$
- ☐  $F_s$
- ☒  $1/(1 - F_p)$
- ☐  $1/F_p$



問題 32

2.5 / 2.5 分數

**A Chunksize of 2:**

- ☐ Breaks your array into 2 even pieces
- ☐ Uses two threads only
- ☒ Deals two for-loop passes to each thread and then goes around to each thread again, etc.
- ☐ Breaks your array into 2 uneven pieces



問題 33

2.5 / 2.5 分數

**The difference between using OpenMP Tasks vs. using OpenMP Sections is that:**

- ☐ Nothing -- they are different words for the same thing
- ☐ Tasks are statically allocated, sections are dynamic
- ☒ Tasks are dynamically allocated, sections are static
- ☐ Sections are deprecated



問題 34

2.5 / 2.5 分數

### The advantage of using the OpenMP *reduction* clause is

- ☐ It is less likely to result in a compiler error
- ☐ No advantage, it is just cleaner code
- ☐ Actually a disadvantage -- it can produce wrong, non-deterministic answers
- ☒ It greatly speeds, and makes thread-safe, reduction operations



問題 35

2.5 / 2.5 分數

### Gustafson's Observation on Amdahl's Law says:

- ☒ When people buy more cores they often do it to process more data, which results in a larger parallel fraction
- ☐ When people buy more cores they often do it to reduce memory contention, which decreases performance
- ☐ Amdahl's Law only applies when you have a number of cores that is less than or equal to 8
- ☐ Amdahl's law was applicable when it was formulated, but doesn't apply now



問題 36

2.5 / 2.5 分數

### Using “default(none)” in an OpenMP #pragma is:

- ☒ A good idea, but not required
- ☐ Required
- ☐ A deprecated feature of an older version of OpenMP
- ☐ A way to possibly increase performance



問題 37

2.5 / 2.5 分數

### A Monte Carlo probability is computed by:

- ☒ Dividing the number of successes by the number of trials
- ☐ Dividing the number of trials by the number of successes
- ☐ Adding the number of successes to the number of trials
- ☐ Subtracting the number of successes from the number of trials



問題 38

2.5 / 2.5 分數

**Intel recently broke the CPU clock speed record by:**

- ☐ Cooling the chip with liquid FluorInert
- ☐ Cooling the chip with four fans
- ☒ Cooling the chip with liquid helium
- ☐ Running the CPU outside during a colder-than-usual winter



問題 39

2.5 / 2.5 分數

**Speedup Efficiency is defined as:**

- ☒  $S_n/n$
- ☐  $F_p/n$
- ☐  $F_p$
- ☐  $n$



問題 40

2.5 / 2.5 分數

**The *observation* that *clock speed* doubles every 2 years:**

- ☐ Is only correct for CPUs, not GPUs
- ☐ Was never actually observed on real systems
- ☐ Has been correct starting in 1965 and is still happening
- ☒ Was the case for a while, but does not apply anymore

測驗分數： 得分：100；總分：100